

Projekat 2 – Big data

Nikola Đorđević 1567

Kafka producer

```
producer = KafkaProducer(bootstrap_servers='kafka:9092',  
                           value_serializer=lambda x: x.encode('utf-8'))  
  
with open('brightkite_small.csv') as file:  
    reader = csv.reader(file)  
    for row in reader:  
        message = ','.join(row)  
        timestamp = str(datetime.now())  
        message_with_timestamp = f"{message},{timestamp}"  
        producer.send('topic3', value=message_with_timestamp)  
  
producer.flush()
```

- Kafka producer čita red po red iz fajla u kome se nalaze podaci, i pročitani red šalje na odgovarajući topic
- Polja jednog reda se šalju odvojena zarezom
- Za svaki poslati red se dodaje i vremenska oznaka kada je poslat

Spark streaming aplikacija

Čitanje i parsiranje podataka sa Kafka topic-a

```
df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka:9092") \
    .option("subscribe", "topic3") \
    .option("startingOffsets", "earliest") \
    .option("maxOffsetsPerTrigger", 1) \
    .load()

parsed_df = df.selectExpr("CAST(value AS STRING)")

split_col = split(parsed_df['value'], ',')
parsed_df = parsed_df.withColumn("user", split_col.getItem(0))
parsed_df = parsed_df.withColumn("check_in_time", split_col.getItem(1))
parsed_df = parsed_df.withColumn("latitude", split_col.getItem(2))
parsed_df = parsed_df.withColumn("longitude", split_col.getItem(3))
parsed_df = parsed_df.withColumn("location_id", split_col.getItem(4))
parsed_df = parsed_df.withColumn("time_spent", split_col.getItem(5))
parsed_df = parsed_df.withColumn("time_stamp", split_col.getItem(6))

parsed_df = parsed_df.where(parsed_df.user != 'user')
parsed_df = parsed_df.withColumn("check_in_time", to_timestamp(parsed_df["check_in_time"], "yyyy-MM-dd'T'HH:mm:ss.SSSXXX"))
parsed_df = parsed_df.withColumn("time_stamp", from_utc_timestamp(parsed_df["time_stamp"], "UTC"))
```

Statistički proračuni za provedeno vreme na određenoj lokaciji, za određenog korisnika

```
# Prosečno, maksimalno i minimalno vreme koje je svaki korisnik proveo na odredjenoj lokaciji
windowed_df_1 = parsed_df.groupby(window(parsed_df.time_stamp, "10 seconds"), parsed_df.user, parsed_df.location_id) \
    .agg(avg(parsed_df.time_spent).alias("avg_time_spent"),
         min(parsed_df.time_spent).alias("min_time_spent"),
         max(parsed_df.time_spent).alias("max_time_spent"))
```

```
def write_to_influxdb(df, epoch_id):
    # Inicijalizacija InfluxDB klijenta
    token = "Xl99hGe7tyPW-wWrgpZRo8vOxA6bK2nR-X3MEoqkigZqnSG1vSqPKOoBmLZdWpWbYKKMKNEfqAAX4FMoKhd5ug=="
    org = "brightkite-org"
    bucket = "brightkite-bucket"
    client = InfluxDBClient(url="http://influxdb:8086", token=token)

    # Kreiranje instance WriteApi klase
    write_api = client.write_api(write_options=SYNCHRONOUS)

    df = df\
        .withColumn("avg_time_spent", col("avg_time_spent").cast("double")) \
        .withColumn("min_time_spent", col("min_time_spent").cast("double")) \
        .withColumn("max_time_spent", col("max_time_spent").cast("double"))

    for row in df.collect():
        point = Point("statistics_3") \
            .tag("window", row.window) \
            .tag("location_id", row.location_id) \
            .tag("user", row.user) \
            .field("avg_time_spent", row.avg_time_spent) \
            .field("min_time_spent", row.min_time_spent) \
            .field("max_time_spent", row.max_time_spent)
        write_api.write(bucket=bucket, org=org, record=point)
```

- Za vremenski prozor od 10 sekundi, podaci se grupišu po korisnicima i lokacijama
- Računa se prosečno, minimalno i maksimalno vreme koje je korisnik proveo na lokaciji

Rezultat u InfluxDB bazi podataka

Filter tables...

_field = avg_time_spent_measurement = statistics_3 location_id = 24b42b2485a1

_field = max_time_spent_measurement = statistics_3 location_id = 24b42b2485a1

_field = min_time_spent_measurement = statistics_3 location_id = 24b42b2485a1

_field = avg_time_spent_measurement = statistics_3 location_id = 259bf5386b61

_field = max_time_spent_measurement = statistics_3 location_id = 259bf5386b61

_start	_stop	_time	_value	_field	_measurement	location_id	user	window
2023-03-14 02:50:...	2023-03-21 02:50:...	2023-03-15 04:30:...	113	avg_time_spent	statistics_3	259bf5386b6111de...	0	Row(start=datetim...

Broj korisnika koji je posetio svaku od lokacija

```
# Broj korisnika koji je posetio svaku od lokacija
windowed_df_2 = parsed_df.groupby(window(parsed_df.time_stamp, "10 seconds"), parsed_df.location_id) \
    .agg(approx_count_distinct(parsed_df.user).alias("distinct_users"))
```

```
def write_to_influxdb_2(df, epoch_id):
    # Inicijalizacija InfluxDB klijenta
    token = "Xl99hGe7tyPW-wWrgpZRo8vOxA6bK2nR-X3MEoqkigZqnSG1vSqPKOoBmLZdWpWbYKKMKNEfqAAX4FMoKhd5ug=="
    org = "brightkite-org"
    bucket = "brightkite-bucket"
    client = InfluxDBClient(url="http://influxdb:8086", token=token)

    # Kreiranje instance WriteApi klase
    write_api = client.write_api(write_options=SYNCHRONOUS)

    for row in df.collect():
        point = Point("dist_users") \
            .tag("window", row.window) \
            .tag("location_id", row.location_id) \
            .field("distinct_users", row.distinct_users)
        write_api.write(bucket=bucket, org=org, record=point)
```

- Određivanje broja korisnika koji je posetio svaku od lokacija za vremenski prozor od 10 sekundi, i upis rezultata u InfluxDB bazu podataka

Rezultat u InfluxDB bazi podataka

Filter tables...

_field = distinct_users _measurement = dist_users location_id = 02aa8df4ed1d7fb

_field = distinct_users _measurement = dist_users location_id = 115b5206a20011d

_field = distinct_users _measurement = dist_users location_id = 1505c4070d66c4

_field = distinct_users _measurement = dist_users location_id = 16d31436ae0711dc

_field = distinct_users _measurement = dist_users location_id = 17363f11b927617c

_start	_stop	_time	_value	_field	_measurement	location_id	window
2023-03-14 03:05:43 ...	2023-03-21 03:05:43 ...	2023-03-15 01:00:00 ...	3	distinct_users	dist_users	115b5206a20011dd9f...	Row(start=datetime.d...

Određivanje top N lokacija i upis u InfluxDB

```
windowed_df_3 = parsed_df.groupby(window(parsed_df.time_stamp, "10 seconds"), parsed_df.location_id) \
    .agg(count("*").alias("num_visits"))
```

```
def write_to_influxdb_3(df, epoch_id):
    N = 2
    # Inicijalizacija InfluxDB klijenta
    token = "Xl99hGe7tyPW-wWrgpZRo8vOxA6bK2nR-X3MEoqkigZqnSG1vSqpKOoBmLZdWpWbYKKMKNEfqAAX4FMoKhd5ug=="
    org = "brightkite-org"
    bucket = "brightkite-bucket"
    client = InfluxDBClient(url="http://influxdb:8086", token=token)

    # Kreiranje instance WriteApi klase
    write_api = client.write_api(write_options=SYNCHRONOUS)

    df = df.orderBy(col("num_visits").desc()).limit(N)

    for row in df.collect():
        point = Point("top_n_locations_v2") \
            .tag("location_id", row.location_id) \
            .field("num_visits", row.num_visits)
        write_api.write(bucket=bucket, org=org, record=point)
```

Na osnovu broja poseta za određenu lokaciju, u vremenskom prozoru od 10 sekundi, vrši se sortiranje lokacija i određivanje top N lokacija, koje se upisuju u InfluxDB bazu

Rezultat u InfluxDB bazi podataka

Table

CUSTOMIZE

Local

SAVE AS

Filter tables...

_field = num_visits _measurement = top_n_locations_v2 location_id = 11da318f0ea

_field = num_visits _measurement = top_n_locations_v2 location_id = 6f3a2db56c

_field = num_visits _measurement = top_n_locations_v2 location_id = 6f5b96170b

_field = num_visits _measurement = top_n_locations_v2 location_id = b3d356765

_field = num_visits _measurement = top_n_locations_v2 location_id = f6f52a75fd8

_start	_stop	_time	_value	_field	_measurement	location_id
2023-03-21 05:19:25 GM...	2023-03-21 05:20:25 GM...	2023-03-21 05:20:04 GM...		1 num_visits	top_n_locations_v2	6f3a2db56d4fa788f72de...
2023-03-21 05:19:25 GM...	2023-03-21 05:20:25 GM...	2023-03-21 05:20:14 GM...		2 num_visits	top_n_locations_v2	6f3a2db56d4fa788f72de...

Skripta za pokretanje aplikacije

```
#!/bin/bash
```

```
spark/bin/spark-submit --master spark://spark-master:7077 --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.1.2 app_2.py
```

Web UI



Application: PySpark_Kafka_Consumer

ID: app-20230321150800-0000
Name: PySpark_Kafka_Consumer
User: root
Cores: Unlimited (16 granted)
Executor Limit: Unlimited (2 granted)
Executor Memory: 1024.0 MiB
Executor Resources:
Submit Date: 2023/03/21 15:08:00
State: RUNNING
[Application Detail UI](#)

▼ Executor Summary (2)

ExecutorID	Worker	Cores	Memory	Resources	State	Logs
1	worker-20230321150626-172.18.0.7-37795	8	1024		RUNNING	stdout stderr
0	worker-20230321150627-172.18.0.8-46755	8	1024		RUNNING	stdout stderr

Web UI

[illegible]

Flink aplikacija

Čitanje i parsiranje podataka sa Kafka topic-a

```
public static DataStream<CheckIn> StreamConsumer(String inputTopic, String server, StreamExecutionEnvironment environment) throws Exception {  
    FlinkKafkaConsumer<String> flinkKafkaConsumer = createStringConsumerForTopic(inputTopic, server);  
    DataStream<String> stringInputStream = environment.addSource(flinkKafkaConsumer);  
  
    return stringInputStream.map(new MapFunction<String, CheckIn>() {  
        private static final long serialVersionUID = -999736771747691234L;  
  
        @Override  
        public CheckIn map(String value) throws Exception {  
            String[] split = value.split(regex:",");  
            return new CheckIn(  
                Integer.parseInt(split[0]),  
                split[1],  
                Double.parseDouble(split[2]),  
                Double.parseDouble(split[3]),  
                split[4],  
                Integer.parseInt(split[5])  
            );  
        }  
    });  
}
```

Statistički proračuni za provedeno vreme na određenoj lokaciji

```
SingleOutputStreamOperator groupedStream = dataStream
    .keyBy(CheckIn::getLocation_id)
    .window(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(10)))
    .aggregate(new MinMaxAggregateFunction());

groupedStream.print();

CassandraSink.addSink(groupedStream)
    .setHost("cassandra")
    .setQuery("INSERT INTO flink.min_max(location_id, max, min) values (?, ?, ?);")
    .build();
```

- Određivanje minimalnog i maksimalnog provedenog vremena na svakoj od lokacija i upis rezultata u Cassandra bazu podataka
- Podaci su grupisani u vremenski (po tipu "Sliding") prozor veličine 10 sekundi i koraka 10 sekundi

Rezultat u Cassandra DB

location_id	max	min
a8d770948720fd154ccfe093a21b73e0	201	201
4dfb8a7526182e70b53ea6a7ee670b5a3d694375	277	277
424eb3dd143292f9e013efa00486c907	354	340
e653dbdc794b11ddb6a0030487eb504	347	347
dd7cd3d264c2d063832db506fba8bf79	217	107
b3d356765cc8a4aa7ac5cd18caafd393	200	3
3a955c5c8336ae8040fa61ceece8a0c0	30	30
6f5b96170b7744af3c7577fa35ed0b8f	188	188
08a35293e09f508494096c1c1b3819edb9df50db	267	8
dd98ca10f19211dd948d003048c10834	81	81
6ccba23d2e036ef1bcff1feb17443d6c7ef8f579	210	14
6f3a2db56d4fa788f72def616f79b7a4	254	20
d2f35ed2724211dea6e6003048c0801e	184	184
36740e782a458ebcab9f14ae3a8a5f19d04d0e41	151	151
828fb63e770c11ddb8450030487eb504	261	261
85cd3b30938d11dd9bf0003048c10834	331	331
e2460b3f5040fc666cd99128a0c133b7	333	333
098e8e9f4b606ce204a76dc504b61697	284	284
b3fe66207ee111dd89eb0030487eb504	352	352

Statistički proračuni za provedeno vreme na određenoj lokaciji

```
SingleOutputStreamOperator meanStream = dataStream
    .keyBy(CheckIn::getLocation_id)
    .window(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(10)))
    .aggregate(new MeanAggregate());

meanStream.print();

CassandraSink.addSink(meanStream)
    .setHost("cassandra")
    .setQuery("INSERT INTO flink.mean(location_id, mean) values (?, ?);")
    .build();
```

- Određivanje prosečnog provedenog vremena na svakoj od lokacija i upis rezultata u Cassandra bazu podataka
- Podaci su grupisani u vremenski (po tipu "Sliding") prozor veličine 10 sekundi i koraka 10 sekundi

Rezultat u Cassandra DB

location_id	mean
a8d770948720fd154ccfe093a21b73e0	201
4dfb8a7526182e70b53ea6a7ee670b5a3d694375	277
424eb3dd143292f9e013efa00486c907	347
e653dbdc794b11ddb6a0030487eb504	347
dd7cd3d264c2d063832db506fba8bf79	155.16667
b3d356765cc8a4aa7ac5cd18caafd393	101.5
3a955c5c8336ae8040fa61ceece8a0c0	30
6f5b96170b7744af3c7577fa35ed0b8f	188
08a35293e09f508494096c1c1b3819edb9df50db	137.5
dd98ca10f19211dd948d003048c10834	81
6ccba23d2e036ef1bcff1feb17443d6c7ef8f579	136.66667
6f3a2db56d4fa788f72def616f79b7a4	83.4
d2f35ed2724211dea6e6003048c0801e	184
36740e782a458ebcab9f14ae3a8a5f19d04d0e41	151
828fb63e770c11ddb8450030487eb504	261
85cd3b30938d11dd9bf0003048c10834	331
e2460b3f5040fc666cd99128a0c133b7	333
098e8e9f4b606ce204a76dc504b61697	284
b3fe66207ee111dd89eb0030487eb504	352
0397002b74d5c5ebddf3ea242554782	325

Broj korisnika koji je posetio svaku od lokacija

```
DataStream
```

- Određivanje broja korisnika koji je posetio svaku od lokacija i upis rezultata u Cassandra bazu podataka
- Podaci su grupisani u vremenski (po tipu "Tumbling") prozor veličine 10 sekundi

Rezultat u Cassandra DB

user_location	counts
64b925364ac71005af756eabb94096e2-0	1
9848afcc62e500a01cf6fbf24b797732f8963683-0	2
a2451a4a10d0061e87ebcd21b422f44b-0	1
7a0f88982aa015062b95e3b4843f9ca2-0	5
e2460b3f5040fc666cd99128a0c133b7-0	1
974b77e179e872596500b2d22c38bf26-0	1
c83ac485e066dbf68b8619c7ab5b2579-0	1
8f060f74f59a02df0993b24964334eea-0	2
d2f35ed2724211dea6e6003048c0801e-0	1
828fb63e770c11ddb8450030487eb504-0	1
115b5206a20011dd9f01003048c10834-0	1
dcc06bf19e775f436c2225be50e14922-0	2
c69bd906c5210caee0719913f836303c-0	1
e63d729e56954aeb23ba669d2c7a2805-0	1
6f5b96170b7744af3c7577fa35ed0b8f-0	1
ee1f60000f61a0000000000000000000-0	1

Top N lokacija

```
DataStream<Tuple2<String, Integer>> topN = dataStream
.map(new MapFunction<CheckIn, String>() {
    @Override
    public String map(CheckIn value) throws Exception {
        return value.getLocation_id();
    }
})
.map(new MapFunction<String, Tuple2<String, Integer>>() {
    @Override
    public Tuple2<String, Integer> map(String value) throws Exception {
        return new Tuple2<String, Integer>(value, 1);
    }
})
.keyBy(new KeySelector<Tuple2<String, Integer>, String>() {
    @Override
    public String getKey(Tuple2<String, Integer> value) throws Exception {
        return value.f0;
    }
})
//.window(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(10)))
.window(TumblingProcessingTimeWindows.of(Time.seconds(10)))
.sum(1)
//.windowAll(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(10)))
.windowAll(TumblingProcessingTimeWindows.of(Time.seconds(10)))
.process(new ProcessAllWindowFunction<Tuple2<String, Integer>, Tuple2<String, Integer>, TimeWindow>() {
    @Override
    public void process(Context context, Iterable<Tuple2<String, Integer>> iterable, Collector<Tuple2<String, Integer>> collector) throws Exception {
        PriorityQueue<Tuple2<String, Integer>> queue = new PriorityQueue<>(Comparator.comparingInt(o -> o.f1));
        for (Tuple2<String, Integer> t : iterable) {
            queue.offer(t);
            if (queue.size() > N) {
                queue.poll();
            }
        }
        List<Tuple2<String, Integer>> topN = new ArrayList<>(queue);
        topN.sort(Comparator.comparingInt(o -> -o.f1));
        for (Tuple2<String, Integer> t : topN) {
            collector.collect(t);
        }
    }
});
```

- Pronalazak Top N lokacija u odnosu na broj poseta
- Podaci su grupisani u vremenski (po tipu "Tumbling") prozor veličine 10 sekundi

Rezultat u Cassandra DB

location_id	counts
dd7cd3d264c2d063832db506fba8bf79	6
6f3a2db56d4fa788f72def616f79b7a4	5
2ef143e12038c870038df53e0478cefc	9
f6f52a75fd80e27e3770cd3a87054f27	7
7a0f88982aa015062b95e3b4843f9ca2	5

Top N lokacija - Cassandra DB

```
CassandraSink.addSink(topN)
    .setHost("cassandra")
    .setQuery("INSERT INTO flink.top_n_locations(location_id, counts) values (?, ?);")
    .build();
```


Web UI

The screenshot displays the Apache Flink Dashboard Web UI. On the left is a dark blue sidebar with navigation links: Overview, Jobs (expanded), Running Jobs, Completed Jobs, Task Managers, Job Manager, and Submit New Job (highlighted in blue). The main content area has a top header with 'Version: 1.16.1', 'Commit: DeadD0d0 @ 1970-01-01T01:00:00+01:00', and 'Message: 0'. Below this is the 'Uploaded Jars' section, which includes a table with one entry: 'Flink-1.0-SNAPSHOT.jar' uploaded on '2023-03-21, 17:52:10' with entry class 'proj.App'. Below the table are input fields for 'proj.App', 'Parallelism', 'Program Arguments', and 'Savepoint Path'. There are also checkboxes for 'Allow Non Restored State' and 'Show Plan', and a 'Submit' button.

Name	Upload Time	Entry Class	
Flink-1.0-SNAPSHOT.jar	2023-03-21, 17:52:10	proj.App	Delete

Below the table, there are input fields for:

- proj.App
- Parallelism
- Program Arguments
- Savepoint Path

There are also checkboxes for:

- Allow Non Restored State
- Show Plan

A 'Submit' button is located at the bottom right of the form.

- Aplikacija se startuje tako što se upload-uje željeni JAR fajl preko web UI - a

Web UI

Apache Flink Dashboard

Overview

Jobs

Running Jobs

Completed Jobs

Task Managers

Job Manager

Submit New Job

Version: 1.16.1 Commit: DeadD0d0 @ 1970-01-01T01:00:00+01:00 Message: 0

Flink Streaming Job [Cancel Job](#)

Job ID	8930f3b1ea96015c409005ad034aa150	Job State	RUNNING 6	Actions	Job Manager Log
Start Time	2023-03-21 17:53:59	Duration	9s		

[Overview](#) [Exceptions](#) [TimeLine](#) [Checkpoints](#) [Configuration](#)

```
graph LR; S["Source: Custom Source -> Map  
-> Map -> Map -> Map  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"] -- "100%" --> T1["TumblingProcessingTimeWindows  
-> Sink: Print to Std. Out. To  
RR: Cassandra Sink  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"]; S -- "100%" --> T2["TumblingProcessingTimeWindows  
-> Sink: Print to Std. Out. To  
RR: Cassandra Sink  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"]; S -- "100%" --> T3["SlidingProcessingTimeWindows  
-> Sink: Print to Std. Out. To  
RR: Cassandra Sink  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"]; S -- "100%" --> T4["SlidingProcessingTimeWindows  
-> Sink: Print to Std. Out. To  
RR: Cassandra Sink  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"]; T1 -- "100%" --> T5["TriggerWindowTumblingProcessingTimeWindows(10000, List  
SinkAndSourceParameters=window  
contents, defaultWindow=roll, so  
window=org.apache.flink.api.co  
mmon.operators.windowing.operators  
.WindowOperator$RollingOperator  
@4465d47f, ProcessingTimeWi  
ndow(1, org.apache.flink.api.co  
mmon.operators.windowing.operators  
.WindowOperator$RollingOperator  
@4465d47f) -> Sink: Print  
To Std. Out. To RR: Cassandra Sink  
Parallelism: 1  
Backgrounded (read): YES  
Keep state: YES"];
```

Web UI

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Tasks
Source: Custom Source -> Map -> (Map -> Map, Map -> Map)	RUNNING	0 B	0	0 B	198	1	2023-03-21 17:...	1
TumblingProcessingTimeWindows -> (Sink: Print to Std. Out, Sink: Ca...	RUNNING	5.15 KB	99	0 B	0	1	2023-03-21 17:...	1
TumblingProcessingTimeWindows	RUNNING	4.96 KB	99	0 B	0	1	2023-03-21 17:...	1
TriggerWindow(TumblingProcessingTimeWindows(10000), ListStateD...	RUNNING	4 B	0	0 B	0	1	2023-03-21 17:...	1
SlidingProcessingTimeWindows -> (Sink: Print to Std. Out, Sink: Cassa...	RUNNING	10.5 KB	99	0 B	0	1	2023-03-21 17:...	1
SlidingProcessingTimeWindows -> (Sink: Print to Std. Out, Sink: Cassa...	RUNNING	10.5 KB	99	0 B	0	1	2023-03-21 17:...	1